

Running Jobs on Comet (a practical guide)

Mary Thomas (mthomas@sdsc.edu)

April 19, 2019

Part 1

Outline

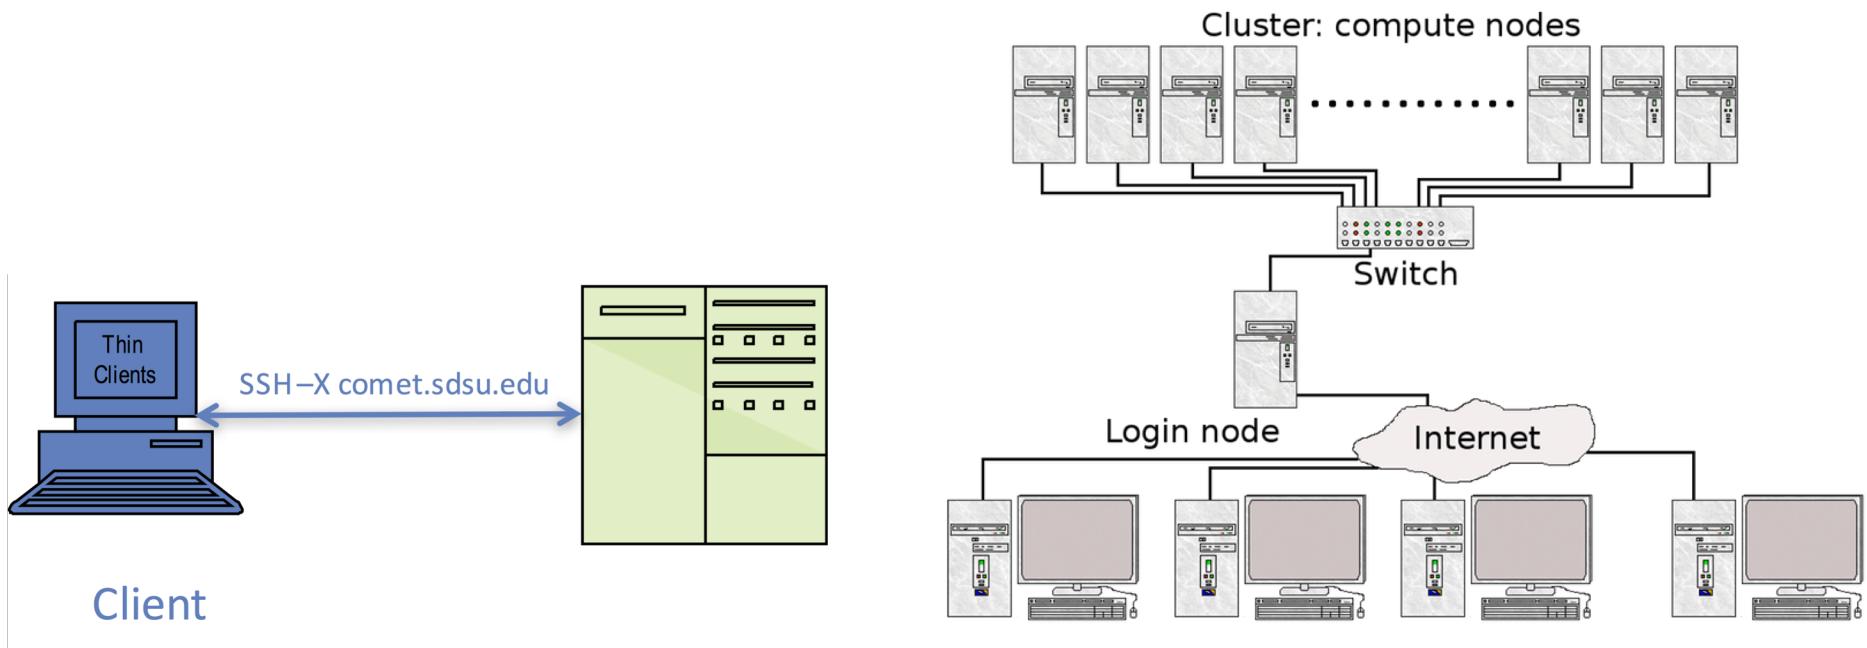
- **Getting Started/Comet System Environment**
- **Comet Overview**
- **Compiling and Linking Code**
- **Running Parallel Jobs**
 - Running OpenMP Jobs
 - Running MPI Jobs
 - Running Hybrid MPI-OpenMP Jobs
 - Running GPU/CUDA Jobs
- **Final Comments**

Getting Started

Basic Information

- Comet User Guide:
 - https://www.sdsc.edu/support/user_guides/comet.html
- Online repo for companion tutorial/webinar information:
 - <https://github.com/sdsc-training/webinars>
 - You must be familiar with running basic Unix commands: see the basic_skills and getting_started links in the webinar directory cited above.
- You must have a comet account in order to access the system. To obtain a trial account:
 - http://www.sdsc.edu/support/user_guides/comet.html#trial_accounts
- More training events listed at SDSC:
 - https://www.sdsc.edu/education_and_training/training.html

System Access: Logging On



- Linux/Mac – Use available ssh clients.
- Putty, Cygwin - ssh clients for windows:
 - <http://www.chiark.greenend.org.uk/~sgtatham/putty/>
- Login hostname for SDSC Comet: comet.sdsc.edu (198.202.113.252)
- Passwordless login tutorial:
 - <https://www.tecmint.com/ssh-passwordless-login-using-ssh-keygen-in-5-easy-steps/>

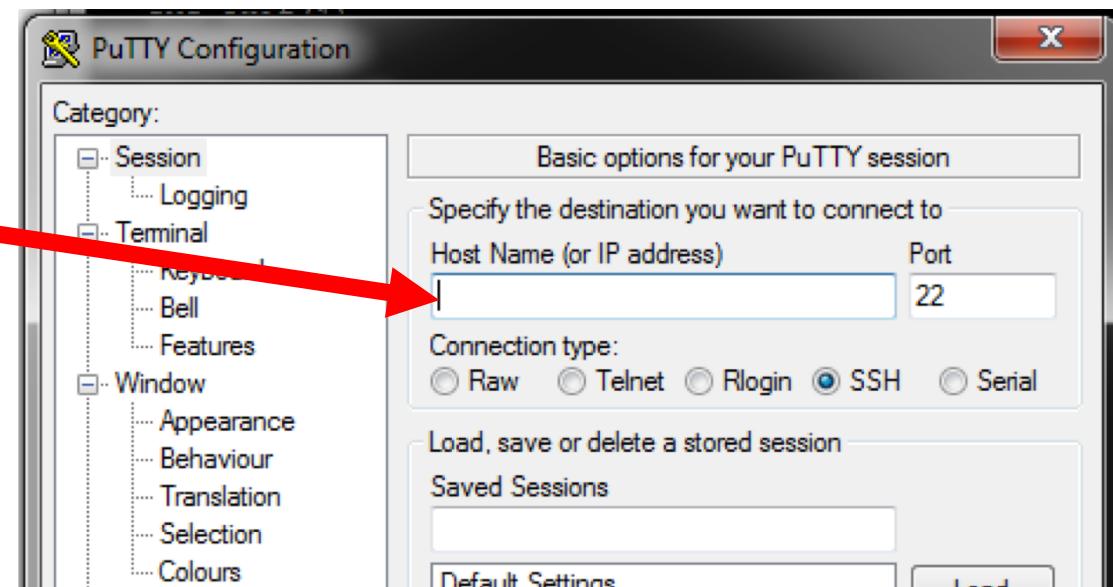
Logging into Comet

Mac/Linux:

`ssh username@comet.sdsc.edu`

Windows (PuTTY):

comet.sdsc.edu



Example of a terminal connection:

```
[$USER@wireless-169-228-105-171:~] ssh comet.sdsc.edu
Warning: No xauth data; using fake authentication data for X11 forwarding.
Last login: Mon Jan  7 15:01:50 2019 from wireless-169-228-105-171.ucsd.edu
Rocks 6.2 (SideWinder)
Profile built 16:45 08-Feb-2016
```

```
Kickstarted 17:27 08-Feb-2016
```

```
WELCOME TO
```



```
*****
```

```
[1] Example Scripts: /share/apps/examples
```

```
[2] Filesystems:
```

- (a) Lustre scratch filesystem : /oasis/scratch/comet/\$USER/temp_project
(Preferred: Scalable large block I/O)
- (b) Compute/GPU node local SSD storage: /scratch/\$USER/\$SLURM_JOBID
(Meta-data intensive jobs, high IOPs)
- (c) Lustre projects filesystem: /oasis/projects/nsf
- (d) /home/\$USER : Only for source files, libraries, binaries.
Do not use for I/O intensive jobs.

```
[3] Comet User Guide: http://www.sdsc.edu/support/user\_guides/comet.html
```

```
*****
```

```
[$USER@comet-ln3:~]
```

Obtaining Tutorial Example Code

- Create a test directory hold the comet example files (e.g. comet-examples)
- Copy the **PHYS244** directory from the/share/apps/examples directory to your 'comet-examples' directory
- This tutorial will focus on examples in bold.

```
[$USER@comet-ln2 ~]$ mkdir comet-examples
[username@comet-ln3 ~]$ cp -r /share/apps/examples/PHYS244/ comet-examples/
[$USER@comet-ln3:~/comet-examples] ls -al PHYS244/
total 230
drwxr-xr-x 16 user use300 16 Aug  5 19:02 .
drwxr-xr-x  5 user use300  6 Aug  5 19:02 ..
drwxr-xr-x  2 user use300  5 Aug  5 19:02 COMPILER_EXAMPLES
drwxr-xr-x  2 user use300 14 Aug  6 00:56 CUDA
drwxr-xr-x  2 user use300 11 Aug  5 19:02 debug
drwxr-xr-x  3 user use300  3 Aug  5 19:02 HADOOP
drwxr-xr-x  2 user use300  6 Aug  6 00:12 HYBRID
drwxr-xr-x  2 user use300  6 Aug  5 19:02 LOCALSCRATCH
drwxr-xr-x  2 user use300  5 Aug  5 19:02 LOCALSCRATCH2
drwxr-xr-x  2 user use300  9 Nov 25 17:29 MKL
drwxr-xr-x  4 user use300  7 Aug  6 09:55 MPI
drwxr-xr-x  2 user use300  8 Aug  5 19:02 OpenACC
drwxr-xr-x  2 user use300  8 Aug  5 23:25 OPENMP
drwxr-xr-x  3 user use300  5 Aug  5 19:02 pytorch
drwxr-xr-x  4 user use300  4 Aug  5 19:02 SPARK
drwxr-xr-x  4 user use300  5 Aug  5 19:02 TensorFlow
```

Comet Overview

HPC for the “long tail of science:”



Img Src taken from:

<https://www.cc.gatech.edu/~echow/ipcc/hpc-course/HPC-networks.pdf>

Stone Soupercomputer: cheapest cost/flop=\$0
~20 MFlops

<https://web.archive.org/web/20031121211117/http://stonesoup.esd.ornl.gov/>



Comet: HPC for the “long tail of science:”

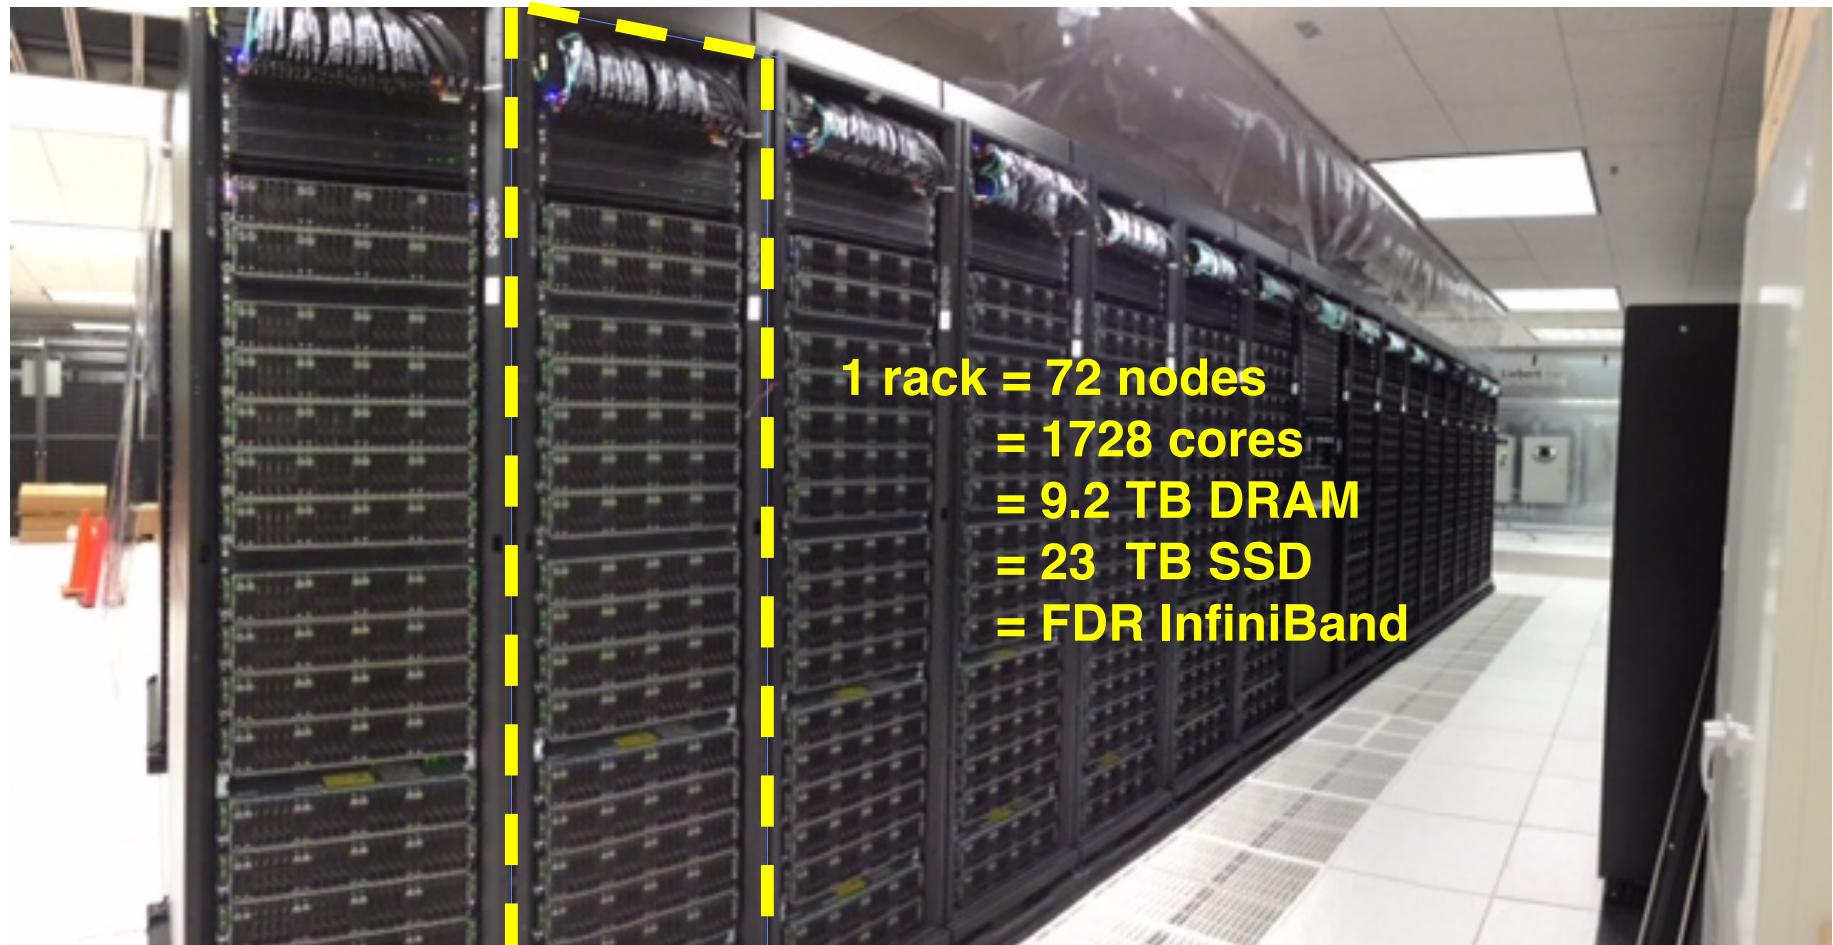
- “Long Tail” - majority of computational research is performed at *modest scale*: large number jobs that run for less than 48 hours, but can be computationally intensive and generate large amounts of data.
- Comet is an NSF-funded system available through the eXtreme Science and Engineering Discovery Environment (XSEDE) program.
- Advanced computing environment: supports science gateways, interactive computing, Jupyter notebooks, containers.



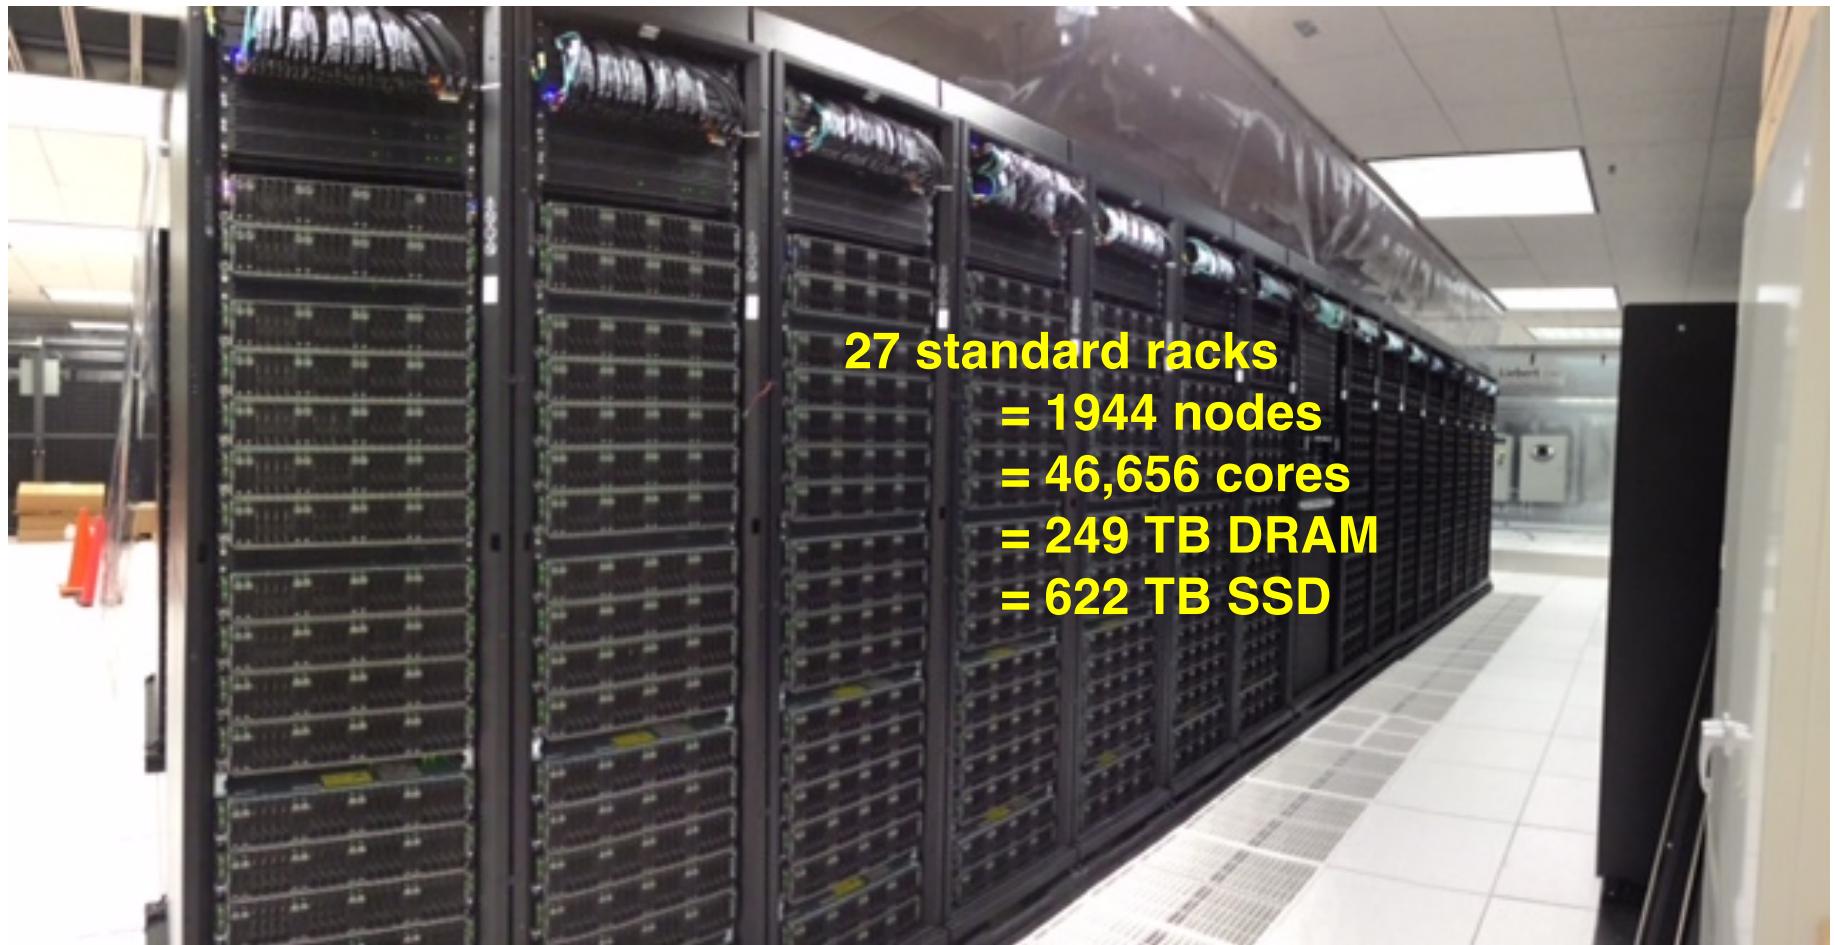
Comet: System Characteristics

- Total peak flops ~2.76 PF
- Dell primary integrator
 - Intel Haswell processors w/ AVX2
 - Mellanox FDR InfiniBand
- 1944 Standard compute nodes (46,656 cores)
 - Dual CPUs, each 12-core, 2.5 GHz
 - 128 GB DDR4 2133 MHz DRAM
 - 2*160GB GB SSDs (local disk)
- 72 GPU nodes
 - 36 nodes same as standard nodes *plus* Two NVIDIA K80 cards, each with dual Kepler3 GPUs
 - 36 nodes with 2 14-core Intel Broadwell CPUs plus 4 NVIDIA P100 GPUs
- 4 large-memory nodes
 - 1.5 TB DDR4 1866 MHz DRAM
 - Four Haswell processors/node
 - 64 cores/node
- Hybrid fat-tree topology
 - FDR (56 Gbps) InfiniBand (bisection)
 - Rack-level (72 nodes, 1,728 cores) full bisection bandwidth
 - 4:1 oversubscription cross-rack
- Performance Storage
 - 7.6 PB, 200 GB/s; Lustre
 - Scratch & Persistent Storage segments
- Durable Storage
 - 6 PB, 100 GB/s; Lustre
 - Automatic backups of critical data
- Home directory storage
- Gateway hosting nodes
- Virtual image repository
- 100 Gbps external connectivity to Internet2 & ESNet

~67 TF supercomputer in a rack



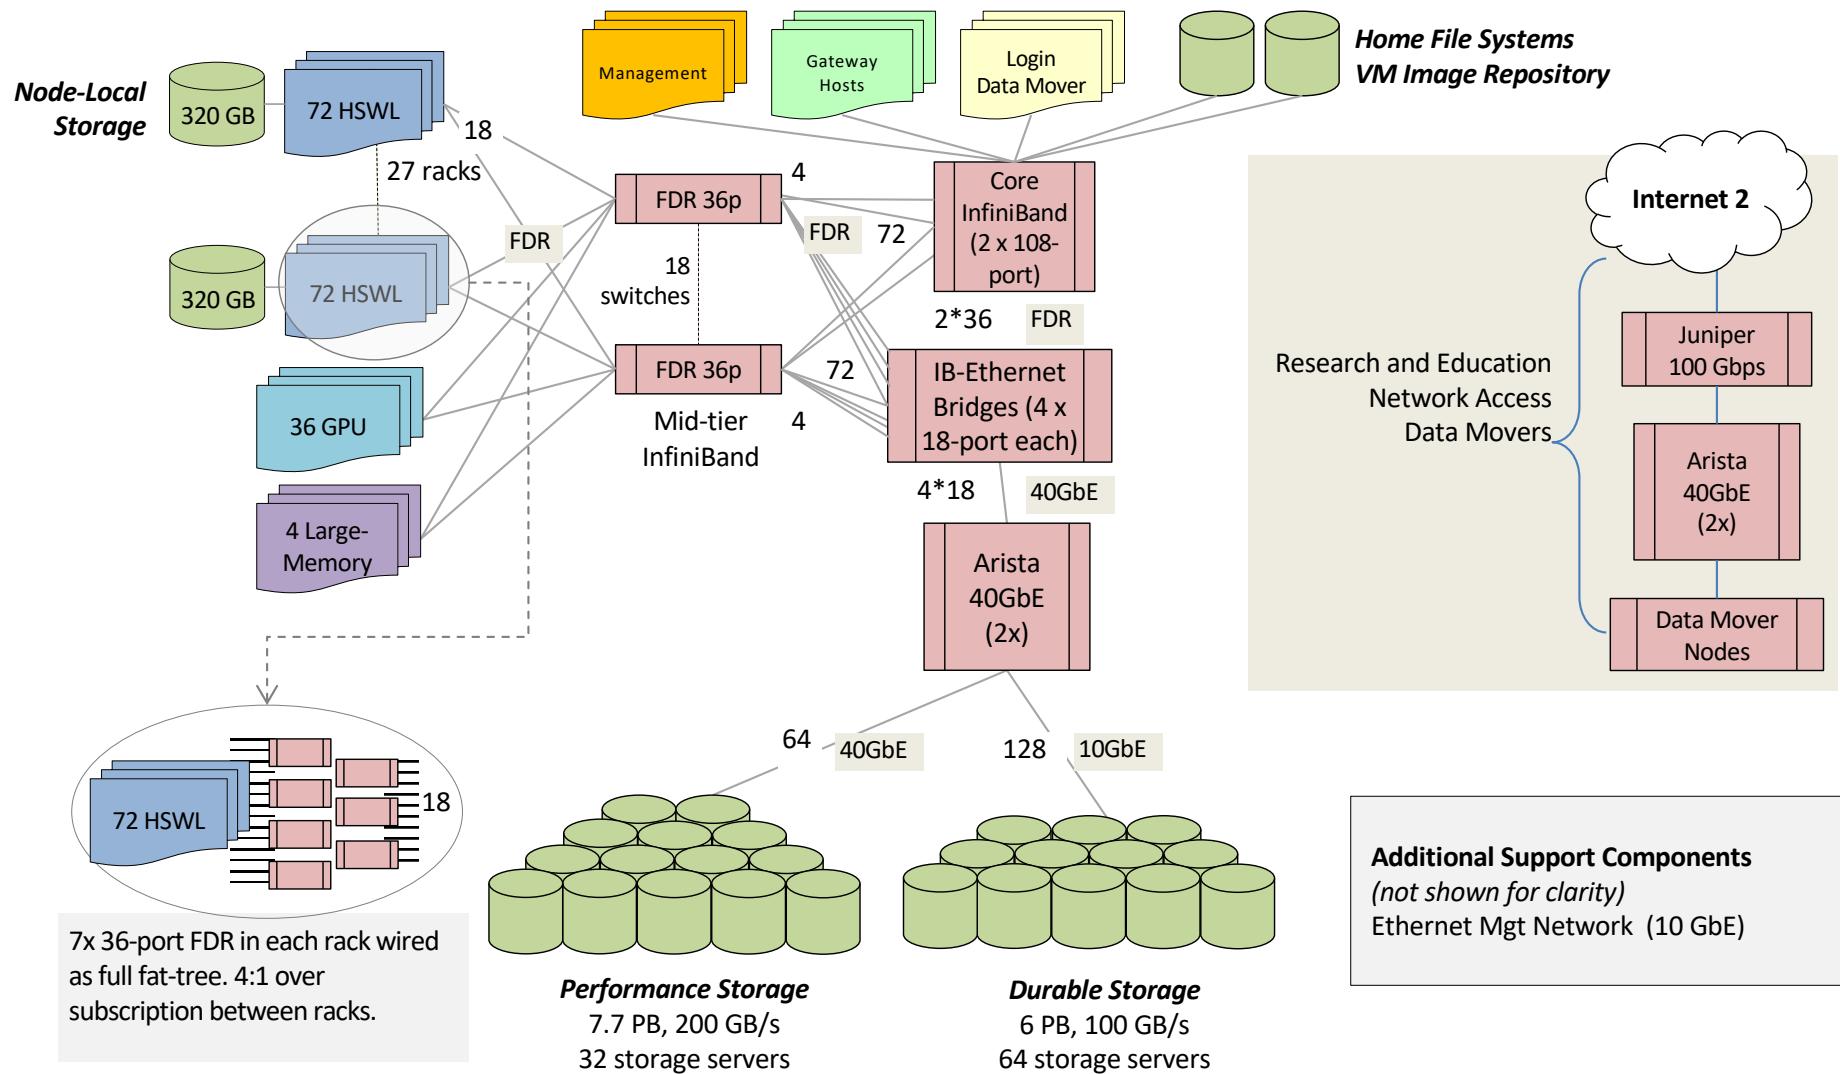
And 27 single-rack supercomputers



27 standard racks
= 1944 nodes
= 46,656 cores
= 249 TB DRAM
= 622 TB SSD

Comet Network Architecture

InfiniBand compute, Ethernet Storage



Comet: Filesystems

- Lustre filesystems – Good for scalable large block I/O
 - Accessible from all compute and GPU nodes.
 - /oasis/scratch/comet - 2.5PB, peak performance: 100GB/s. Good location for storing large scale scratch data during a job.
 - /oasis/projects/nsf - 2.5PB, peak performance: 100 GB/s. Long term storage.
 - ***Not good for large # of small files or small block I/O.*** Please. Don't.
- SSD filesystems
 - /scratch local to each native compute node – 210GB on regular compute nodes, 285GB on GPU, large memory nodes, 1.4TB on selected compute nodes.
 - SSD location is good for writing small files and temporary scratch files. Purged at the end of a job.
- Home directories (/home/\$USER)
 - Source trees, binaries, and small input files.
 - ***Not good for large scale I/O.***

Comet File Systems

Path	Purpose	User Access Limits	Lifetime
\$HOME	NFS storage; Source code, important files	100 GB	Backed-up
/oasis/scratch/comet/ \$USER/temp_project	Global/Parallel Lustre FS; temp storage for distributed access	500 GB	No backup
/oasis/projects/nsf	Global/Parallel Lustre FS; project storage	~2.5 PB total	Backed-up
/scratch/\$USER/\$SL URM_JOB_ID	Local SSD on batch job node fast per-node access	210 GB per compute node, 286GB on GPU, Large memory nodes	Purged after job ends

Managing the Environment with Modules

Comet: System Environment

- Modules are used to manage environment for users.
- Default environment:

\$ module li

Currently Loaded Module files:

1) intel/2013_sp1.2.144 2) mvapich2_ib/2.1 3) gnutools/2.69

- Listing available modules:

\$ module av

----- /opt/modulefiles/mpi/.intel -----

intelmpi/2016.3.210(default) mvapich2_ib/2.1(default)

mvapich2_gdr/2.1(default) openmpi_ib/1.8.4(default)

mvapich2_gdr/2.2

----- /opt/modulefiles/applications/.intel -----

atlas/3.10.2(default) lapack/3.6.0(default) scalapack/2.0.2(default)

boost/1.55.0(default) mxml/2.9(default) slepc/3.6.2(default)

...

...

Modules: Managing the User Environment

A few popular commands

Command	Description
module list	List the modules that are currently loaded
module avail	List the modules that are available
module display <module_name>	Show the environment variables used by and how they are affected
module show <module_name>	Same as display
module unload	Remove from the environment
module load	Load into the environment
module swap	Replace with in the environment

Module Command Examples

- Default environment: list, li

```
[$USER@comet-ln3:~/comet-examples] module li
Currently Loaded Module files: 1) intel/2013_sp1.2.144 2) mvapich2_ib/2.1      3) gnutools/2.69
```

- List available modules: available, avail, av

```
[$USER@comet-ln3:~/comet-examples] module av

----- /opt/modulefiles/applications/.intel -----
-----
atlas/3.10.2(default)      hdf5/1.8.14(default)      papi/5.4.1(default)      sundials/2.6.2(default)
boost/1.55.0(default)       ipm/2.0.3(default)       parmetis/4.0.3(default)  superlu/4.2(default)
fftw/2.1.5                 lapack/3.6.0(default)    pdt/3.20(default)        tau/2.23(default)
fftw/3.3.4(default)        mxml/2.9(default)        petsc/3.6.3(default)
trilinos/11.12.1(default) 
gsl/1.16                   netcdf/3.6.2           scalapack/2.0.2(default)
gsl/2.1(default)           netcdf/4.3.2(default)    slepc/3.6.2(default)
hdf4/2.11(default)         p3dfft/2.7.4(default)   sprng/2.0b(default)

$ module av ----- /opt/modulefiles/mpi/.intel -----
intelmpi/2016.3.210(default) mvapich2_ib/2.1(default)
mvapich2_gdr/2.1(default)   openmpi_ib/1.8.4(default) mvapich2_gdr/2.2
----- /opt/modulefiles/applications/.intel -----
atlas/3.10.2(default)      lapack/3.6.0(default)    scalapack/2.0.2(default)
boost/1.55.0(default)       mxml/2.9(default)        slepc/3.6.2(default)

... MORE....
```

Module Command Examples

- Load a module, and show what it does

```
[$USER@comet-ln3:~/comet-examples] env
HOSTNAME=comet-ln3.sdsc.edu
IPPREROOT=/opt/intel/composer_xe_2013_sp1.2.144/ipp
INTEL_LICENSE_FILE=/opt/intel/composer_xe_2013_sp1.2.144/licenses:/opt/intel/licenses:/root/intel/licenses
TERM=xterm-256color
SHELL=/bin/bash
HISTSIZE=5000
GDBSERVER_MIC=/opt/intel/composer_xe_2013_sp1.2.144/debugger/gdb/target/mic/bin/gdbserver
SSH_CLIENT=169.228.105.171 58704 22
[SNIP]
HOME=/home/user
ROLLSREROOT=/opt/rocks/share-devel/src/roll
MPIHOME=/opt/mvapich2/intel/ib
FFTWHOME=/opt/fftw/3.3.4/intel/mvapich2_ib
SDSCHOME=/opt/sdsc
PYTHONPATH=/opt/sdsc/lib
LOGNAME=user
QTLIB=/usr/lib64/qt-3.3/lib
CVS_RSH=ssh
SSH_CONNECTION=169.228.105.171 58704 198.202.113.252 22
MODULESHOME=/usr/share/Modules
MKL_ROOT=/opt/intel/composer_xe_2013_sp1.2.144/mkl
LESSOPEN=||/usr/bin/lesspipe.sh %
INFOPATH=/opt/intel/composer_xe_2013_sp1.2.144/debugger/gdb/intel64/share/info:/opt/intel/composer_xe_2013_sp1.2.144/debugger/gdb/intel64_mic/share/info/
DISPLAY=localhost:42.0
INCLUDE=/opt/intel/composer_xe_2013_sp1.2.144/mkl/include
INTELHOME=/opt/intel/composer_xe_2013_sp1.2.144
G_BROKEN_FILERAMES=1
BASH_FUNC_module()=() { eval `/usr/bin/modulecmd bash $*` }
_=~/bin/env
```

Module: check Environment

Once you have loaded the modules, you can check the system variables that are available for you to use.

```
[$USER@comet-ln3:~/comet-examples] module load fftw/3.3.4
[$USER@comet-ln3:~/comet-examples]
[$USER@comet-ln3:~/comet-examples] module li
Currently Loaded Modulefiles:
 1) intel/2013_sp1.2.144  2) mvapich2_ib/2.1          3) gnutools/2.69      4)
fftw/3.3.4
[$USER@comet-ln3:~/comet-examples] module show fftw/3.3.4
-----
/opt/modulefiles/applications/.intel/fftw/3.3.4:

module-whatis          fftw
module-whatis          Version: 3.3.4
module-whatis          Description: fftw
module-whatis          Compiler: intel
module-whatis          MPI Flavors: mvapich2_ib openmpi_ib
setenv                 FFTWHOME /opt/fftw/3.3.4/intel/mvapich2_ib
prepend-path            PATH /opt/fftw/3.3.4/intel/mvapich2_ib/bin
prepend-path            LD_LIBRARY_PATH /opt/fftw/3.3.4/intel/mvapich2_ib/lib
prepend-path            LIBPATH /opt/fftw/3.3.4/intel/mvapich2_ib/lib
-----
```

Using Script to load modules

Control and guarantee the current working environment. In order for the commands run inside a script (child shell) to change the parent shell, you must use the **source** command.

```
[comet-ln3:~] source ./loadgpuenv.sh  
[comet-ln3:~] module list  
Currently Loaded Modulefiles:  
 1) gnutools/2.69  2) cuda/7.0
```

```
[comet-ln3:~] which nvcc  
/usr/local/cuda-7.0/bin/nvcc  
[mthomas@comet-ln3:~] which mpirun  
/usr/bin/which: no mpirun in (/opt-gnu/gcc/bin:/usr/local/bin.....)
```

```
[comet-ln3:~] source loadintelenv.sh  
[comet-ln3:~] module list  
Currently Loaded Modulefiles:
```

```
 1) gnutools/2.69  2) intel/2013_sp1.2.144 3) mvapich2_ib/2.1
```

```
[mthomas@comet-ln3:~] which mpirun  
/opt/mvapich2/intel/ib/bin/mpirun  
[mthomas@comet-ln3:~] which nvcc  
/usr/bin/which: no nvcc in (/opt-gnu/gcc/bin:/usr/local/bin.....)
```

```
comet-ln3:~] cat loadgpuenv.sh  
#!/bin/bash  
module purge  
module load gnutools  
module load cuda
```

```
[comet-ln3:~] cat loadintelenv.sh  
module purge  
module load gnutools  
module load intel mvapich2_ib
```

Module: command not found

- Sometimes encountered when switching from one shell to another or attempting to run the module command from within a shell script or batch job.
- Module command may not be inherited to the shell
- To keep this from happening, execute the following command:
 - command line (interactive shells)
 - `source /etc/profile.d/modules.sh`
 - OR add to your shell script (including Slurm batch scripts)

Compiling & Linking

Compiling & Linking: Topics

- Supported Compiler Types
- Intel Compiling
- PGI Compiling
- GNU Compiling
- GPU Compiling

Supported Compiler Types

- Comet compute nodes support several parallel programming models:
 - **MPI**: Default Intel Compiler: `intel/2013_sp1.2.144`;
 - Versions 2015.2.164 and 2016.3.210 available.
 - Other options: `openmpi_ib/1.8.4` (and 1.10.2), Intel MPI, `mvapich2_ib/2.1`
 - `mvapich2_gdr`: GPU direct enabled version
 - **OpenMP & Pthreads**:
 - All compilers (GNU, Intel, PGI) have OpenMP flags.
 - **GPU** nodes: support CUDA, OpenACC.
 - **Hybrid** modes are possible (see examples below).

Suggested Compilers

- Default/Suggested Compilers to used based on programming model and languages:

	Serial	MPI	OpenMP	MPI + OpenMP
Fortran	ifort	mpif90	ifort -openmp	mpif90 -openmp
C	icc	mpicc	icc -openmp	mpicc -openmp
C++	icpc	mpicxx	icpc -openmp	mpicxx -openmp

- In this tutorial, we include hands-on examples that cover many of the cases in the table:
 - (1) MPI
 - (2) OpenMP
 - (3) HYBRID
 - (4) Local scratch

Using the Intel Compilers

- Intel compilers and MVAPICH2 MPI implementation will be loaded by default.
- If you have modified your environment, you can reload by executing the module purge & load commands at the Linux prompt, or placing the load command in your startup file (`~/.cshrc` or `~/.bashrc`)

```
[$USER@comet-ln2:~] module purge
[$USER@comet-ln2:~] module list
No Modulefiles Currently Loaded.
[$USER@comet-ln2:~] module load gnutools
[$USER@comet-ln2:~] module list
Currently Loaded Modulefiles:
 1) gnutools/2.69
[$USER@comet-ln2:~] module load intel mvapich2_ib
[$USER@comet-ln2:~] module list
Currently Loaded Modulefiles:
 1) gnutools/2.69    2) intel/2013_sp1.2.144  3) mvapich2_ib/2.1
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] which mpicc
/opt/mvapich2/intel(ib/bin)/mpicc
```

```
[comet-ln3:~] source loadintelenv.sh
[comet-ln3:~] module list
Currently Loaded Modulefiles:
 1) gnutools/2.69 2) intel/2013_sp1.2.144 3) mvapich2_ib/2.1
[mthomas@comet-ln3:~] which mpirun
/opt/mvapich2/intel(ib/bin)/mpirun
```

Using the Intel Compilers

- For Intel Advanced Vector Extensions (AVX2) support, compile with the `-xHOST` option.
 - https://en.wikipedia.org/wiki/Advanced_Vector_Extensions (128/256bit SIMD, Vector ops (MPI broadcast, gather, ...))
 - Note that `-xHOST` alone does not enable aggressive optimization, so compilation with `-O3` is also suggested.
 - The `-fast` flag invokes `-xHOST`, but should be avoided since it also turns on interprocedural optimization (`-ipo`), which may cause problems in some instances.
- Intel Math Kernel Lib (MKL) libraries are available as part of the "intel" modules on Comet.
 - Once this module is loaded, the environment variable `MKL_ROOT` points to the location of the mkl libraries.
 - The MKL link advisor can be used to ascertain the link line (change the `MKL_ROOT` aspect appropriately).

Using the Intel Compilers

- In the example below, we are working with the HPC examples that can be found in PHYS244/MKL :

```
[$USER@comet-14-01:~]cd comet-examples/PHYS244/MKL
[$USER@comet-14-01:~/comet-examples/PHYS244/MKL] pwd
/home/user/comet-examples/PHYS244/MKL
[$USER@comet-14-01:~/comet-examples/PHYS244/MKL] ls -al
total 25991
drwxr-xr-x  2 user use300      9 Nov 25 17:20 .
drwxr-xr-x 16 user use300     16 Aug  5 19:02 ..
-rw-r--r--  1 user use300    325 Aug  5 19:02 compile.txt
-rw-r--r--  1 user use300   6380 Aug  5 19:02 pdpttr.c
-rwxr-xr-x  1 user use300 44825440 Nov 25 16:55 pdpttr.exe
-rw-r--r--  1 user use300    188 Nov 25 16:57 scalapack.20294236.comet-07-27.out
-rw-r--r--  1 user use300    376 Aug  5 19:02 scalapack.sb
```

Using the Intel Compilers

The file `compile.txt` contains the full command to compile the `pdpttr.c` program statically linking 64 bit scalapack libraries on Comet:

```
[$USER@comet-14-01:~/comet-examples/PHYS244/MKL] cat compile.txt
mpicc -o pdpttr.exe pdpttr.c -I${MKL_ROOT}/include
${MKL_ROOT}/lib/intel64/libmkl_scalapack_lp64.a -Wl,--start-group
${MKL_ROOT}/lib/intel64/libmkl_intel_lp64.a
${MKL_ROOT}/lib/intel64/libmkl_core.a
${MKL_ROOT}/lib/intel64/libmkl_sequential.a -Wl,--end-group
${MKL_ROOT}/lib/intel64/libmkl_blacs_intelmpi_lp64.a -lpthread -lm
```

Verify your environment, then compile the command:

```
[$USER@comet-14-01 :~/comet-examples/PHYS244/MKL] source ~/loadintelenv.sh
[$USER@comet-14-01 :~/comet-examples/PHYS244/MKL] module list
Currently Loaded Modulefiles:
 1) gnutools/2.69          2) intel/2013_sp1.2.144   3) mvapich2_ib/2.1
[$USER@comet-14-01:~/comet-examples/PHYS244/MKL] mpicc -o pdpttr.exe pdpttr.c -I${MKL_ROOT}/include
${MKL_ROOT}/lib/intel64/libmkl_scalapack_lp64.a -Wl,--start-group
${MKL_ROOT}/lib/intel64/libmkl_intel_lp64.a ${MKL_ROOT}/lib/intel64/libmkl_core.a
${MKL_ROOT}/lib/intel64/libmkl_sequential.a -Wl,--end-group
${MKL_ROOT}/lib/intel64/libmkl_blacs_intelmpi_lp64.a -lpthread -lm
```

For more information on the Intel compilers run: [ifort | icc | icpc] -help

Using the PGI Compilers

- PGI (formerly The Portland Group, Inc.), was a company that produced a set of commercially available Fortran, C and C++ compilers for high-performance computing systems.
- It is now owned by NVIDIA.
- PGI compilers can be loaded by executing the following commands at the Linux prompt or placing in your startup file (~/.cshrc or ~/.bashrc).
- For AVX support, compile with -fast

```
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] module purge
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] module load gnutools
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] module load pgi
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] module load mvapich2_ib
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] module list
Currently Loaded Modulefiles:
    1) gnutools/2.69      2) pgi/17.5          3) mvapich2_ib/2.1
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] which mpicc
/opt/mvapich2/pgi(ib/bin/mpicc
```

- For more information on the PGI compilers run: man [pgf90 | pgcc | pgCC]

Recommended PGI Compilers

	Serial	MPI	OpenMP	MPI+OpenMP
Fortran	pgf90	mpif90	pgf90 -mp	mpif90 -mp
C	pgcc	mpicc	pgcc -mp	mpicc -mp
C++	pgCC	mpicxx	pgCC -mp	mpicxx -mp

- PGI supports the following high-level languages:
 - Fortran 77, 90/95/2003, 2008 (partial)
 - High Performance Fortran (HPF)
 - ANSI C99 with K&R extensions
 - ANSI/ISO C++
 - CUDA Fortran
 - OpenCL
 - OpenACC
 - OpenMP

Using the GNU Compilers

- The GNU compilers can be loaded by executing the following commands at the Linux prompt or placing in your startup files (~/.cshrc or ~/.bashrc)

```
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] module purge
Unloading compiler-dependent module gnutools/2.69
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] module load gnutools
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] module load gnu
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] module load openmpi_ib
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL]
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL] which mpicc
/opt/openmpi/gnu(ib/bin/mpicc
[$USER@comet-ln2:~/comet-examples/PHYS244/MKL]
```

- For AVX support, compile with -mavx.
- Note that AVX support is only available in version 4.7 or later, so it is necessary to explicitly load the gnu/4.9.2 module until such time that it becomes the default.

Using the GNU Compilers

Table of recommended GNU compilers:

	Serial	MPI	OpenMP	MPI+OpenMP
Fortran	gfortran	mpif90	gfortran -fopenmp	mpif90 -fopenmp
C	gcc	mpicc	gcc -fopenmp	mpicc -fopenmp
C++	g++	mpicxx	g++ -fopenmp	mpicxx -fopenmp

Running Jobs On Comet

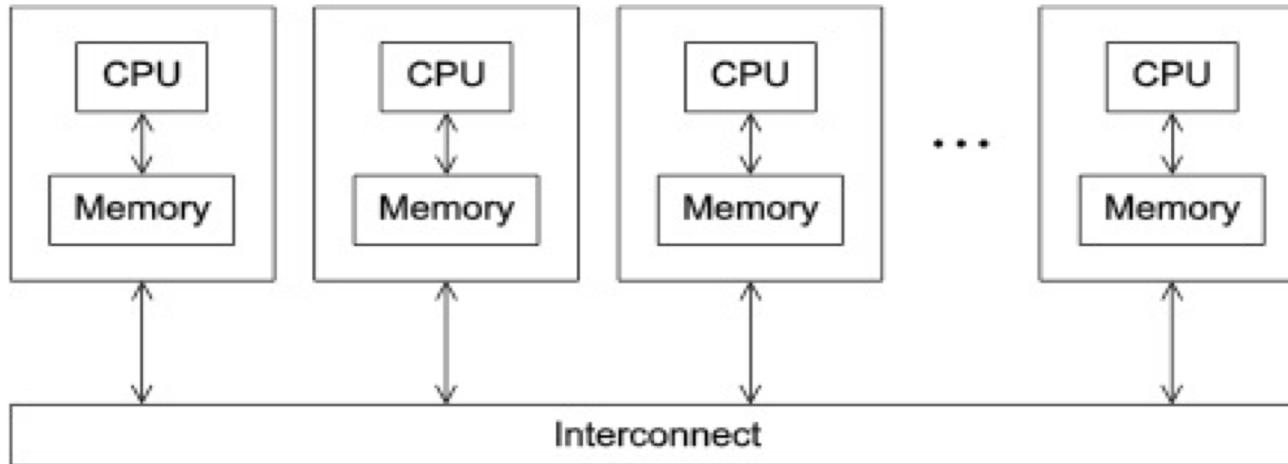
Factors Impacting Job Execution

- **Parallel Models:**
 - Impacts language used, libraries, performance.
- **How you choose to run the job:**
 - Command line execution
 - Batch/queuing System -- Comet uses the Simple Linux Utility for Resource Management (SLURM):
 - Batch Queue
 - Interactive jobs
- **Data I/O choices (topic of upcoming Webinar):**
 - https://www.sdsc.edu/education_and_training/training.html

Parallel Models: Memory

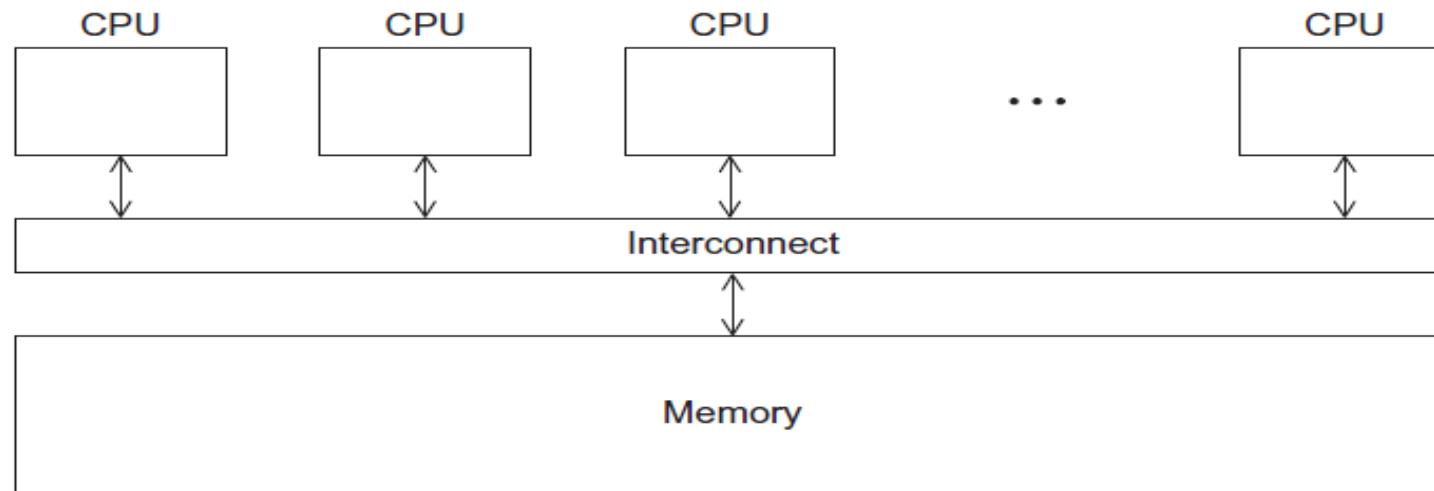
- **Distributed Memory**
- **Shared Memory**
- **Implemented in several languages:**
 - FORTRAN, C, Python, OOPs (sort-of)
- **Large number of libraries and API's**
- **Adds to compilation/linking complexity**

Distributed Memory



- Programs that run **asynchronously**, pass messages for communication and coordination between resources.
- Examples include: SOA-based systems, massively multiplayer online games, peer-to-peer apps.
- Different types of implementations for the message passing mechanism: HTTP, RPC-like connectors, message queues
- HPC historically uses the **Message Passing Interface (MPI)**

Parallel Models: Shared Memory



- **CPUs all share same localized memory (SHMEM);**
 - Coordination and communication between tasks via interprocessor communication (IPC) or virtual memory mappings.
- **May use: uniform or non-uniform memory access (UMA or NUMA); cache-only memory architecture (COMA).**
- **Most common HPC API's for using SHMEM:**
 - Portable Operating System Interface (POSIX); Open Multi-Processing (OpenMP) designed for parallel computing – best for multi-core computing.

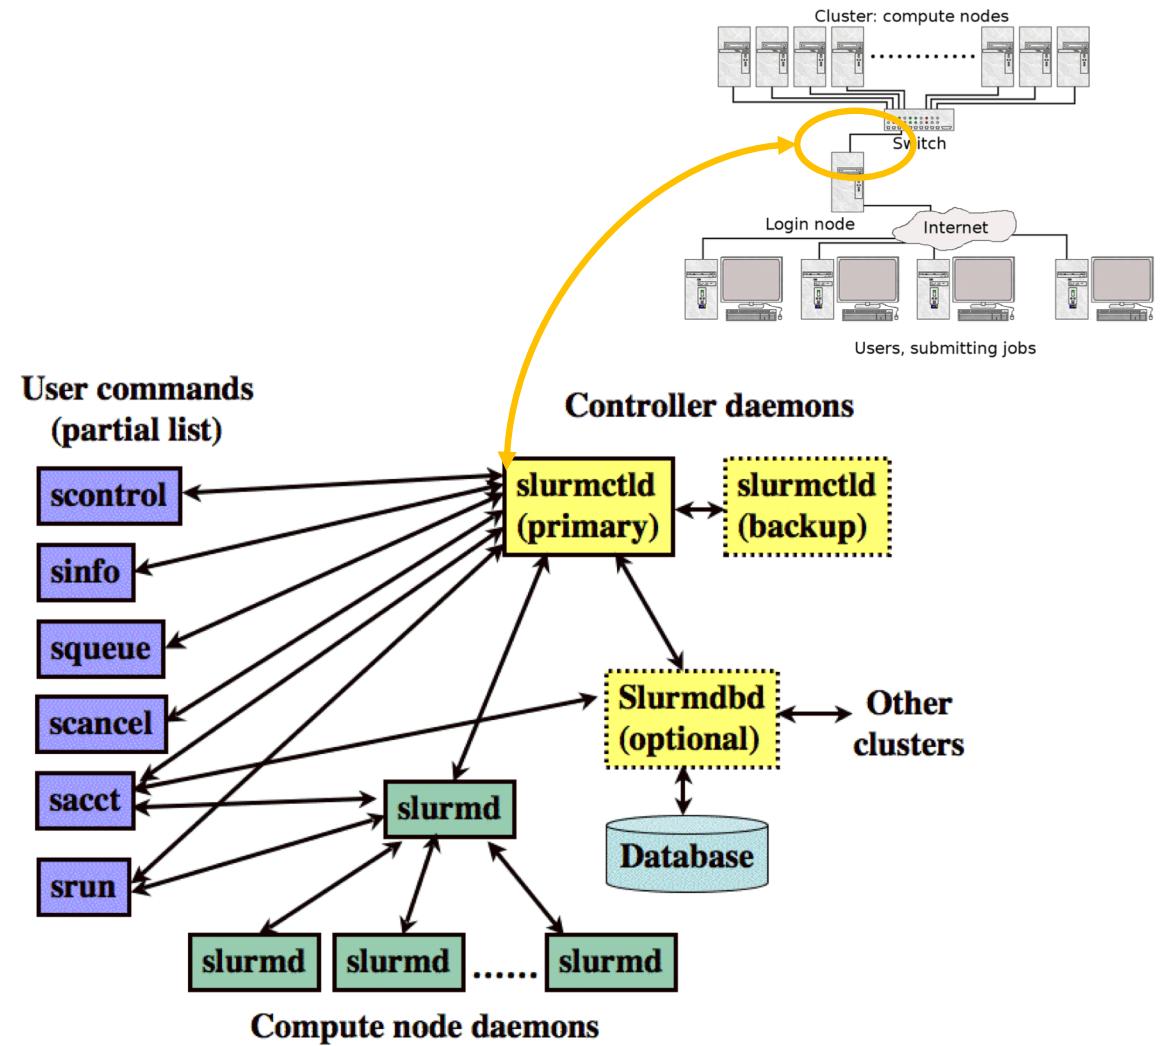
Running Jobs on Comet

- **Important note:** Do not run on the login nodes - even for simple tests.
- All job runs must be via the Slurm scheduling infrastructure.
 - **Interactive Jobs:** Use **srun** command to obtain nodes for ‘live’ interactive access:
`srun --pty --nodes=1 --ntasks-per-node=24 -p debug -t 00:30:00 --wait 0 /bin/bash`
 - **Batch Jobs:** Submit batch scripts from the login nodes. Can choose:
 - Partition (details on upcoming slide)
 - Time limit for the run (maximum of 48 hours)
 - Number of nodes, tasks per node
 - Memory requirements (if any)
 - Job name, output file location
 - Email info, configuration

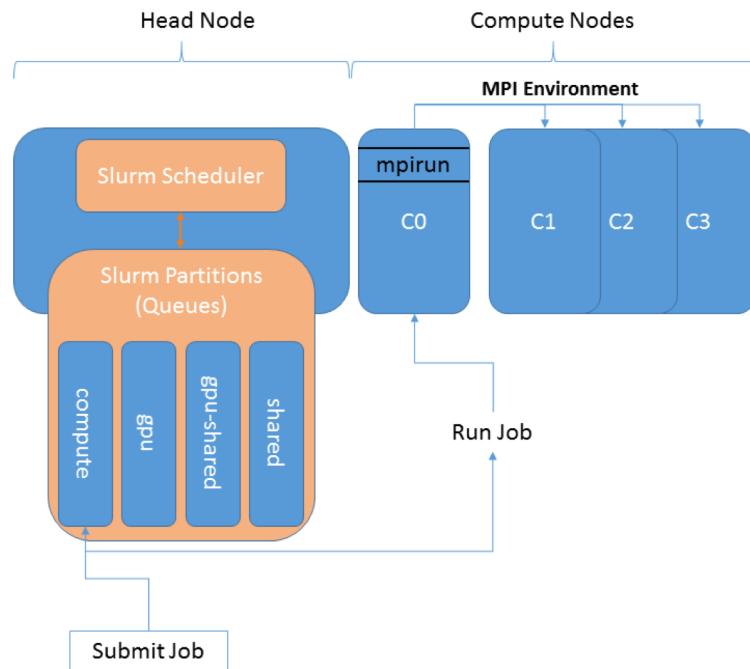
Slurm Resource Manager

Simple Linux Utility for Resource Management

- “Glue” for parallel computer to schedule and execute jobs
- Role: Allocate resources within a cluster
 - Nodes (unique IP address)
 - Interconnect/switches
 - Generic resources (e.g. GPUs)
 - Launch and otherwise manage jobs
- Functionality:
 - Prioritize queue(s) of jobs;
 - decide when and where to start jobs;
 - terminate job when done;
 - Appropriate resources;
 - manage accounts for jobs



Slurm Partitions on Comet



Specified using -p option in **sbatch** script. For example:
#SBATCH -p gpu

Queue Name	Max Walltime	Max Nodes	Comments
compute	48 hrs	72	Used for access to regular compute nodes
gpu	48 hrs	4	Used for access to the GPU nodes
gpu-shared	48 hrs	1	Used for shared access to a partial GPU node
shared	48 hrs	1	Single-node jobs using fewer than 24 cores
large-shared	48 hrs	1	Single-node jobs using large memory up to 1.45 TB
debug	30 mins	2	Used for access to debug nodes

Common Slurm Commands

- Submit jobs using the **sbatch** command:

```
$ sbatch mycode-slurm.sb
```

Submitted batch job 8718049

- Check job status using the **squeue** command:

```
$ squeue -u $USER
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
8718049	compute	mycode	user	PD	0:00	1	(Priority)

- Once the job is running:

```
$ squeue -u $USER
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
8718064	debug	mycode	user	R	0:02	1	comet-14-01

For Fun:

- **Join the UCSD Supercomputing Club:**
 - <http://supercomputingclub.ucsd.edu/>
 - <https://training.sdsc.edu/scc-training-schedule-dev>
 - Free pie....
 - Raspberry PI^3 event tomorrow, 4/12/19 @ 3pm
- **Check out the Student Cluster Competition Activity @ SDSC:**
 - <https://training.sdsc.edu/scc> (kickoff on 4/12/19 @1pm)
 - Working with the new ARM architecture (RISC)
 - Seeking a few grad students interested in mentoring ☺
- **Take a tour of SDSC! SC Club on 4/19/19**

References

- **Comet User Guide**
 - https://www.sdsc.edu/support/user_guides/comet.html#compiling
- **SDSC Training Resources**
 - https://www.sdsc.edu/education_and_training/training.html
 - <https://github.com/sdsc-training/webinars>
 - Comet shared apps/examples; can be found in
 - /share/apps
- **XSEDE Training Resources**
 - <https://www.xsede.org/for-users/training>
 - <https://cvw.cac.cornell.edu/comet/>